



Debugging (1)

- writing programs is hard
- we need strategies and tools for when programs go wrong
 - or things to guide us to working programs

An example program to debug

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(int argc, char *argv[]) {
    char message[128];
    size_t message_len = 256;
    char timestamp[128];
    time_t t;
    struct tm *tmp;
    FILE *file = fopen(argv[1], "a+");

    printf("Type your log: ");
    getline(&message, &message_len, stdin);

    t = time(NULL);
    tmp = localtime(&t);
    strftime(timestamp, 256, "%C", tmp);

    fprintf(file, "%s: %s\n", timestamp, message);
    return 0;
}
```

- when we compile we get various warnings

and when we run it:

```
./journal <<<"Hello World!"  
Segmentation fault (core dumped)
```

gdb (GNU debugger)

- follows the chain of function calls
- to consistently view the state of the stack
- to view the computer's internals as your program is running

You can set breakpoints with the b command:

```
b program.c:14
```

- the program name followed by the line you wish to set breakpoint at, in this case line 14
 - you will find out what line the program is crashing on by using gdb
- so breakpoints allow you to pause the program and examine its current state
 - e.g. the state of the registers
 - memory
 - other critical data
- this all helps pin point where the program is breaking

Lets attempt to debug it using gdb:

```

# gdb ./journal
Reading symbols from ./journal...
(No debugging symbols found in ./journal)
(gdb) run <<<"hello"
Starting program: /home/joseph/Repos/Talks/COMS10012-Software-Tools/Debugging/journal <<<"hello"
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib64/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.
__vfprintf_internal (s=0x0, format=0x402026 "%s: %s\n", ap=ap@entry=0x7fffffffde50, mode_flags=mode_flags@entry=0) at vfpr
722^I ORIENT;
(gdb) bt
#0 __vfprintf_internal (s=0x0, format=0x402026 "%s: %s\n",
    ap=ap@entry=0x7fffffffde50, mode_flags=mode_flags@entry=0)
    at vfprintf-internal.c:722
#1 0x00007ffff7e2360a in __fprintf (stream=<optimized out>,
    format=<optimized out>) at fprintf.c:32
#2 0x00000000040125f in main ()

```

- we run gdb with: `gdb ./journal`
- the we run it with our arguments: `(gdb) run <<<"hello"` - the <<< was mentioned in pipes, it feeds a string in as input to the program
- in the rest of the output we see that the program crashes as soon as it receives: `SIGSEGV, Segmentation fault`

```

Program received signal SIGSEGV, Segmentation fault.
__vfprintf_internal (s=0x0, format=0x402026 "%s: %s\n", ap=ap@entry=0x7fffffffde50, mode_flags=mode_flags@entry=0) at vfpr
722^I ORIENT;
(gdb) bt
#0 __vfprintf_internal (s=0x0, format=0x402026 "%s: %s\n",
    ap=ap@entry=0x7fffffffde50, mode_flags=mode_flags@entry=0)
    at vfprintf-internal.c:722
#1 0x00007ffff7e2360a in __fprintf (stream=<optimized out>,
    format=<optimized out>) at fprintf.c:32
#2 0x00000000040125f in main ()

```

- in the additional information it says

getting more information

- if we compile with some extra flags for debugging: `cc -Og -g journal.c -o journal`
 - `-g` enables debug symbols
 - `-Og` says “when you are compiling, dont optimise for speed or safety, instead compile to make the program as **easy to debug as possible**”

we then run the program using gdb again:

```

cc -Og -g journal.c -o journal
gdb ./journal
(gdb) run <<"hello"
Starting program: /home/joseph/Repos/Talks/COMS10012-Software-Tools/Debugging/journal <<"hello"
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib64/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.
__memcpy_avx_unaligned_erms () at ../sysdeps/x86_64/multiarch/memmove-vec-unaligned-erms.S:333
Downloading 0.01 MB source file /usr/src/debug/glibc-2.36.9000-19.fc38.x86_64/string/..../sysdeps/x86_64/multiarch/memmove-v
333^I^I^Imovl^I%ecx, -4(%rdi, %rdx)
(gdb) bt
#0 __memcpy_avx_unaligned_erms ()
    at ../sysdeps/x86_64/multiarch/memmove-vec-unaligned-erms.S:333
#1 0x00007ffffe496ac in __GI_getdelim (
    lineptr=lineptr@entry=0x7fffffffdf0, n=n@entry=0x7fffffffde8,
    delimiter=delimiter@entry=10, fp=0x7ffff7fa5aa0 <_IO_2_1_stdin.>
) at iogetdelim.c:111
#2 0x00007ffff7e237d1 in __getline (lineptr=lineptr@entry=0x7fffffffdf0,
    n=n@entry=0x7fffffffde8, stream=<optimized out>) at getline.c:28
#3 0x000000004011d6 in main (argc=<optimized out>, argv=<optimized out>)
    at journal.c:14

```



- the #0 #1 #2 #3 give additional info
- #3 tells us that the program is crashing on line 14 of journal.c

Now within gdb, lets examine what is happening at line 14...

Setting a breakpoint

- use the `(gdb) b` command to set a breakpoint at a given line e.g.

```
(gdb) b journal.c: 14
```

- when we run we hit the breakpoint before triggering the seg fault
- we can see below we are at:

```
14^I getline(&message, &message_len, stdin);
```

```

Breakpoint 2, main (argc=<optimized out>, argv=<optimized out>) at journal.c:14
14^I getline(&message, &message_len, stdin);
(gdb) inspect message
$3 = "@\000\000\000\000\000\000\000\000\200", '\000' <repeats 14 times>, "\006\000\000\000\216\000\000\000\f\000\000\000\b
(gdb) inspect message_len
$4 = 256
(gdb) d
Delete all breakpoints? (y or n) y
(gdb)

```

- so it is crashing when reading in this message

lets use the `inspect` command for a further look:

```
(gdb) inspect message
```

so why is it seg faulting?

If in doubt... read the manual

In man 3 getline:

*getline() reads an entire line from stream, storing the address of the buffer containing the text into *lineptr. The buffer is null-terminated and includes the newline character, if one was found.*

*If *lineptr is set to NULL before the call, then getline() will allocate a buffer for storing the line. This buffer should be freed by the user program even if getline() failed.*

*Alternatively, before calling getline(), *lineptr can contain a pointer to a malloc(3)-allocated buffer *n bytes in size. If the buffer is not large enough to hold the line, getline() resizes it with realloc(3), updating *lineptr and *n as necessary.*

Well we're passing a statically allocated buffer... lets fix that.

- when we wrote the initial program code we didnt give message a size using `malloc` but instead a static length of 128:

```
char message[128];
```

lets fix the source code and run again

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(int argc, char *argv[]) {
    char *message = NULL;
    size_t message_len;
    char timestamp[128];
    time_t t;
    struct tm *tmp;
    FILE *file = fopen(argv[1], "a+");

    printf("Type your log: ");
    getline(&message, &message_len, stdin);

    t = time(NULL);
    tmp = localtime(&t);
    strftime(timestamp, 256, "%C", tmp);

    fprintf(file, "%s: %s\n", timestamp, message);
    return 0;
}

cc -g -Og journal2.c -o journal2
```

when running we still get a seg fault:

```
$ ./journal2 <<<"hello"
Segmentation fault (core dumped)
```

so lets run it using gdb again:

```
# gdb ./journal2
(gdb) run <<<"hello"
Starting program: /home/joseph/Repos/Talks/COMS10012-Software-Tools/Debugging/journal2 <<<"hell
```

now the program is crashing at line 20: `at journal.c: 20`

- this displays every system call your program makes when it is run
- the left side are all the functions that happen when your program is being **loaded into memory**
- the right side is what is happening **whilst the program is running**

Strace lets you use **regular expressions** to filter what system calls you view

- could also use **grep**

Ltrace

- another tool is **ltrace**
 - ltrace is **Linux specific** tracing tool
 - it traces **library calls**
 - where strace traces calls to the operating system (system calls) ltrace traces calls to third party libraries

Running on our program:

```
$ ltrace ./journal3 documents/log.txt <<<hello
fopen("documents/log.txt", "a+")
printf("Type your log: ")
getline(0x7ffffebcc0fc8, 0x7ffffebcc0fc0, 0x7f4bfcf40aa0, 0)
time(nil)
localtime(0x7ffffebcc0f38)
strftime("20", 256, "%C", 0x7f4bfcf47640)
fprintf(nil, "%s: %s\n", "20", "hello\n" <no return ...>
--- SIGSEGV (Segmentation fault) ---
+++ killed by SIGSEGV +++
```

I	= nil
	= 15
	= 6
	= 1674045150
	= 0x7f4bfcf47640
	= 2

- we see that fopen is returning nil = **nil** when it should be returning a pointer to the open file
- the **fprintf(nil...)** says that a null pointer is getting through to that fprintf on line 20, indicating where our program may be going wrong

Using strace with the `-e` flag which limits the output of strace

```
$ strace -e openat ./journal3 documents/log.txt <<<hello
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/lib64/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "documents/log.txt", O_RDWR|O_CREAT|O_APPEND, 0666) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/localtime", O_RDONLY|O_CLOEXEC) = 3
--- SIGSEGV {si_signo=SIGSEGV, si_code=SEGV_MAPERR, si_addr=0xc0} ---
+++ killed by SIGSEGV (core dumped) +++
Segmentation fault (core dumped)
```

- we can see that the `(No such file or directory)` is likely causing the issues, i.e. the file you are attempting to open does not currently exist

Lets print some error messages using perror from the `errno.h` C library:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <errno.h>

int main(int argc, char *argv[]) {
    char *message = NULL;
    size_t message_len;
    char timestamp[128];
    time_t t;
    struct tm *tmp;

    if (argc < 2) { printf("Usage %s path/to/log\n", argv[0]); exit(1); }
    FILE *file = fopen(argv[1], "a+"); /* line 11 */
    if (file == NULL) {
        perror("Failed to open log");
        exit(2);
    }

    printf("Type your log: ");
    getline(&message, &message_len, stdin);

    t = time(NULL);
    tmp = localtime(&t);
    strftime(timestamp, 256, "%C", tmp);

    fprintf(file, "%s: %s\n", timestamp, message); /* line 20 */
    return 0;
}
```

Now when we run...

```
$ ./journal4 <<hello
Usage ./journal4 path/to/log

$ ./journal4 documents/log.txt <<hello
Failed to open log: No such file or directory

$ ./journal4 /etc/passwd <<hello
Failed to open log: Permission denied

$ ./journal4 /dev/stdout
Type your log: hello
20: hello
```

- now we're getting some helpful error messages

Valgrind

- tool for detecting memory leaks
- tells you if every memory allocation has a matching free
- running it with the previous program:

Thinking back to when we fixed up getline... it said it would allocate the memory for the line
► ...did we ever free it?

```
$ valgrind ./journal4 /dev/stdout <<<hello
=36111= Memcheck, a memory error detector
=36111= Copyright (C) 2002-2022, and GNU GPL'd, by Julian Seward et al.
=36111= Using Valgrind-3.20.0 and LibVEX; rerun with -h for copyright info
=36111= Command: ./journal4 /dev/stdout
=36111=
20: hello
```

```
Type your log: =36111=
=36111= HEAP SUMMARY:
=36111=     in use at exit: 592 bytes in 2 blocks
=36111= total heap usage: 13 allocs, 11 frees, 13,684 bytes allocated
=36111=
=36111= LEAK SUMMARY:
=36111=     definitely lost: 120 bytes in 1 blocks
=36111=     indirectly lost: 0 bytes in 0 blocks
=36111=     possibly lost: 0 bytes in 0 blocks
=36111=     still reachable: 472 bytes in 1 blocks
=36111=     suppressed: 0 bytes in 0 blocks
```



- we know that we are leaking memory: `definitely lost: 120 bytes...`

Best practice for aiding debugging is **defensive programming**

- **never ignore compiler warnings**
 - use werror flag
- **always check assumptions**
- **always check function return codes**
- use MORE warning flags if possible

e.g. for C programs:

```
-Wall -Wextra --std=c11 -pedantic
```

 these make the compiler incredibly picky about your code

There are also more tools called **linters** which can be even more picky during compilation

- aka **static analysis tools**

- they give you more warnings than normal

BPF tools

- Linux has a new instrumentation framework called eBPF
 - it lets you get lots of detail about what programs are doing
 - **HIGHLY** Linux specific